

AIM

Implementation of Linear and Logistic Regression on Real-World Datasets

THEORY

Regression

Regression is a supervised machine learning technique used to model the relationship between independent variables (features) and a dependent variable (target). It helps in predicting unknown values based on known data.

There are mainly two types of regression used in this experiment:

◆ **Linear Regression**

Linear Regression is used when the output variable is **continuous**.

It assumes a linear relationship between input variables and output.

Mathematical Model:

$$y = mx + c \quad y = mx + c$$

For multiple variables:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n \quad y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

Where:

- $y \rightarrow$ predicted value

- $x \rightarrow$ input features
- $\beta \rightarrow$ coefficients
- $\beta_0 \rightarrow$ intercept

In this experiment:

Target variable: **BMI**

Type: Continuous

Hence, Linear Regression is suitable.

Working of Linear Regression:

1. Initializes weights.
 2. Fits best straight line minimizing error.
 3. Uses Mean Squared Error to optimize parameters.
 4. Predicts BMI for unseen data.
-

Evaluation Metrics:

- MAE (Mean Absolute Error)
- MSE (Mean Squared Error)
- RMSE (Root Mean Squared Error)
- R² Score (Goodness of fit)

Lower MAE/MSE and higher R² indicate better model performance.

◆ Logistic Regression

Logistic Regression is used for **binary classification problems**.

Instead of predicting continuous values, it predicts probabilities between 0 and 1 using the **Sigmoid function**.

Sigmoid Function:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Output $> 0.5 \rightarrow$ Class 1

Output $< 0.5 \rightarrow$ Class 0

In this experiment:

Target variable: **Stroke**

0 → No Stroke

1 → Stroke

Working of Logistic Regression:

1. Applies linear equation.
 2. Passes result through sigmoid.
 3. Converts probability into binary output.
 4. Uses cross-entropy loss for optimization.
-

Evaluation Metrics:

- Accuracy
- Confusion Matrix
- Precision
- Recall
- F1-Score

These metrics help analyze classification performance.



Dataset Description

The healthcare stroke dataset contains patient health records including:

- Age
- Gender
- Hypertension
- Heart disease
- BMI
- Smoking status

Linear Regression predicts BMI, while Logistic Regression predicts stroke occurrence.

Categorical values are encoded and missing BMI values are handled using mean imputation.

⚠ Limitations

Linear Regression:

- Assumes linearity
- Sensitive to outliers
- Cannot model complex patterns

Logistic Regression:

- Assumes linear decision boundary

- Struggles with imbalanced datasets
- Limited for non-linear relationship

CODE

```
# =====#
# STEP 1: Import Libraries
# =====#

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
from sklearn.preprocessing import LabelEncoder

# =====#
# STEP 2: Load Dataset
# =====#

df = pd.read_csv("healthcare-dataset-stroke-data.csv")

print(df.head())

# =====#
# STEP 3: Data Cleaning
# =====#
```

```
# Drop ID column
df.drop("id", axis=1, inplace=True)

# Fill missing BMI with mean
df["bmi"].fillna(df["bmi"].mean(), inplace=True)

# =====
# STEP 4: Encode Categorical Columns
# =====

le = LabelEncoder()

cat_cols =
["gender", "ever_married", "work_type", "Residence_type", "smoking_status"]

for col in cat_cols:
    df[col] = le.fit_transform(df[col])

print("\nAfter Encoding:")
print(df.head())

# =====
# PART A - LINEAR REGRESSION (Predict BMI)
# =====

print("\n===== LINEAR REGRESSION =====")

X = df.drop("bmi", axis=1)
y = df["bmi"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

lr = LinearRegression()
lr.fit(X_train, y_train)
```

```
y_pred_lr = lr.predict(X_test)

print("MAE:", mean_absolute_error(y_test, y_pred_lr))
print("MSE:", mean_squared_error(y_test, y_pred_lr))
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred_lr)))
print("R2 Score:", r2_score(y_test, y_pred_lr))

plt.scatter(y_test, y_pred_lr)
plt.xlabel("Actual BMI")
plt.ylabel("Predicted BMI")
plt.title("Linear Regression: BMI Prediction")
plt.show()

# =====
# PART B - LOGISTIC REGRESSION (Predict Stroke)
# =====

print("\n===== LOGISTIC REGRESSION =====")

X = df.drop("stroke", axis=1)
y = df["stroke"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

log_model = LogisticRegression(max_iter=1000)
log_model.fit(X_train, y_train)

y_pred_log = log_model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred_log))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred_log))
print("\nClassification Report:\n", classification_report(y_test,
y_pred_log))

sns.heatmap(confusion_matrix(y_test, y_pred_log), annot=True, fmt="d")
plt.xlabel("Predicted")
```

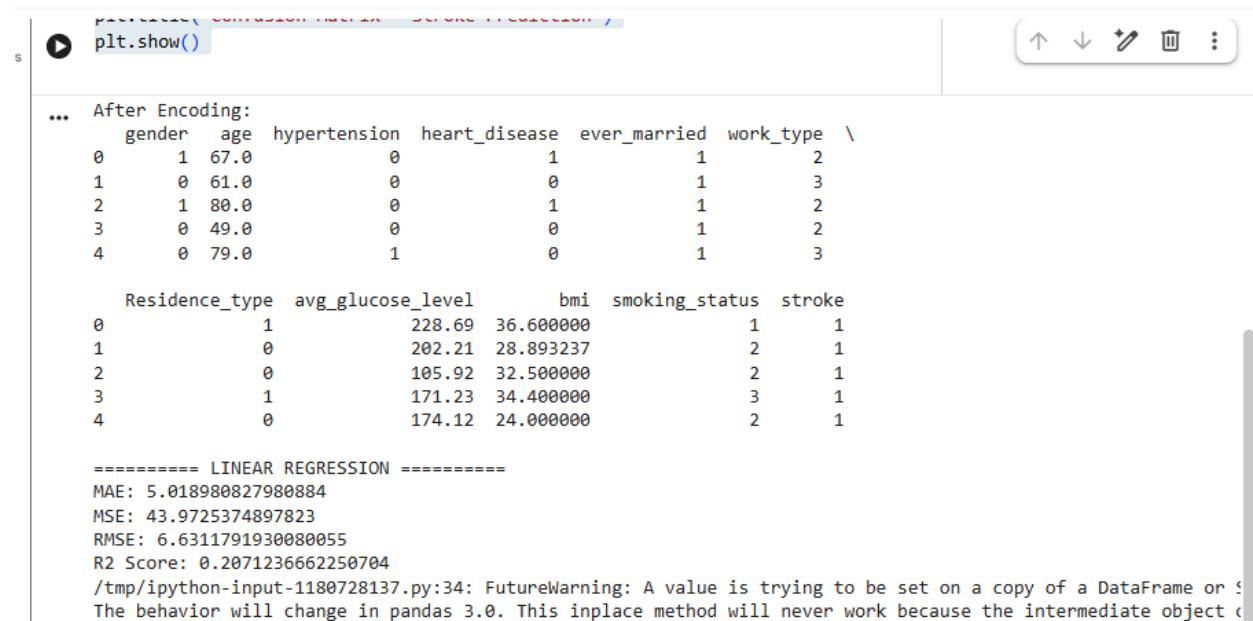
```
plt.ylabel("Actual")
plt.title("Confusion Matrix - Stroke Prediction")
plt.show()
```

OUTPUT

```
...      id  gender  age  hypertension  heart_disease  ever_married \
0    9046   Male  67.0          0            1        Yes
1   51676 Female  61.0          0            0        Yes
2   31112   Male  80.0          0            1        Yes
3   60182 Female  49.0          0            0        Yes
4   1665  Female  79.0          1            0        Yes

      work_type Residence_type  avg_glucose_level    bmi  smoking_status \
0       Private           Urban            228.69  36.6  formerly smoked
1  Self-employed         Rural            202.21    NaN  never smoked
2       Private           Rural            105.92  32.5  never smoked
3       Private           Urban            171.23  34.4      smokes
4  Self-employed         Rural            174.12  24.0  never smoked

      stroke
0      1
1      1
2      1
3      1
4      1
```



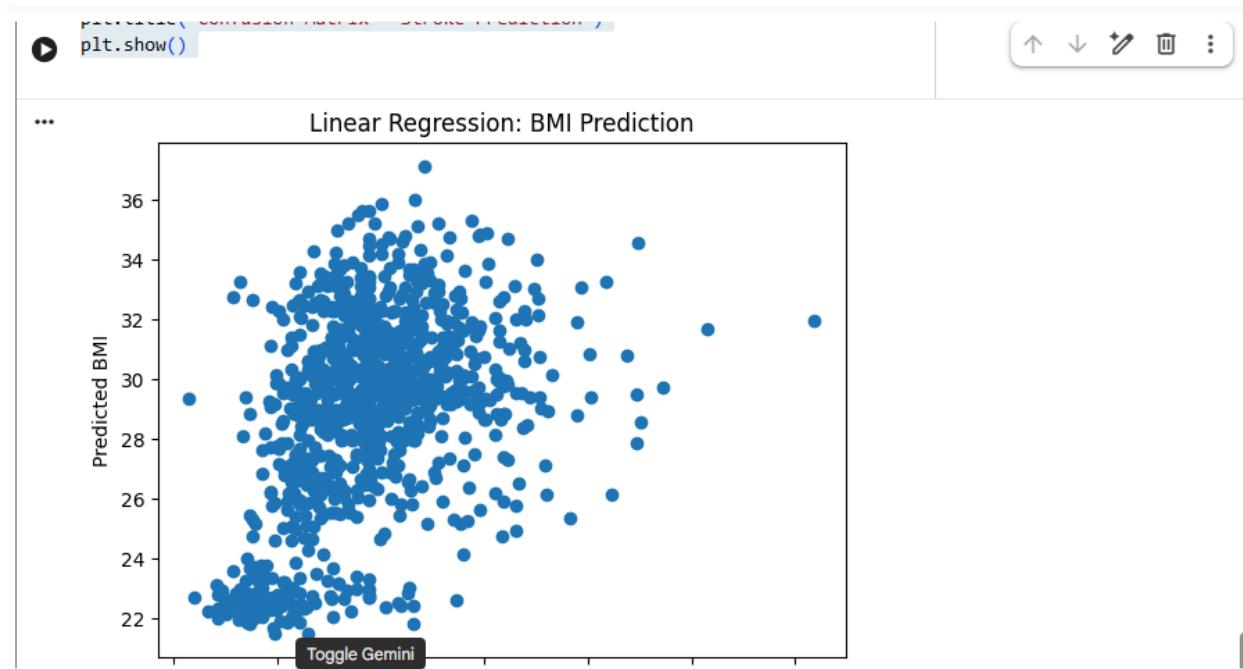
The screenshot shows a Jupyter Notebook interface with the following content:

```
s  plt.show()

... After Encoding:
   gender  age  hypertension  heart_disease  ever_married  work_type \
0      1  67.0          0            1            1            2
1      0  61.0          0            0            1            3
2      1  80.0          0            1            1            2
3      0  49.0          0            0            1            2
4      0  79.0          1            0            1            3

      Residence_type  avg_glucose_level    bmi  smoking_status  stroke
0                  1            228.69  36.600000      1            1
1                  0            202.21  28.893237      2            1
2                  0            105.92  32.500000      2            1
3                  1            171.23  34.400000      3            1
4                  0            174.12  24.000000      2            1

===== LINEAR REGRESSION =====
MAE: 5.018980827980884
MSE: 43.9725374897823
RMSE: 6.6311791930080055
R2 Score: 0.2071236662250704
/tmp/ipython-input-1180728137.py:34: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object c
```



```
===== LOGISTIC REGRESSION =====  
** Accuracy: 0.9393346379647749
```

Confusion Matrix:

```
[[960  0]  
 [ 62  0]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.94	1.00	0.97	960
1	0.00	0.00	0.00	62
accuracy			0.94	1022
macro avg	0.47	0.50	0.48	1022
weighted avg	0.88	0.94	0.91	1022

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Pre  
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Pre  
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Pre  
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Confusion Matrix - Stroke Prediction

